

# Reinforcement learning

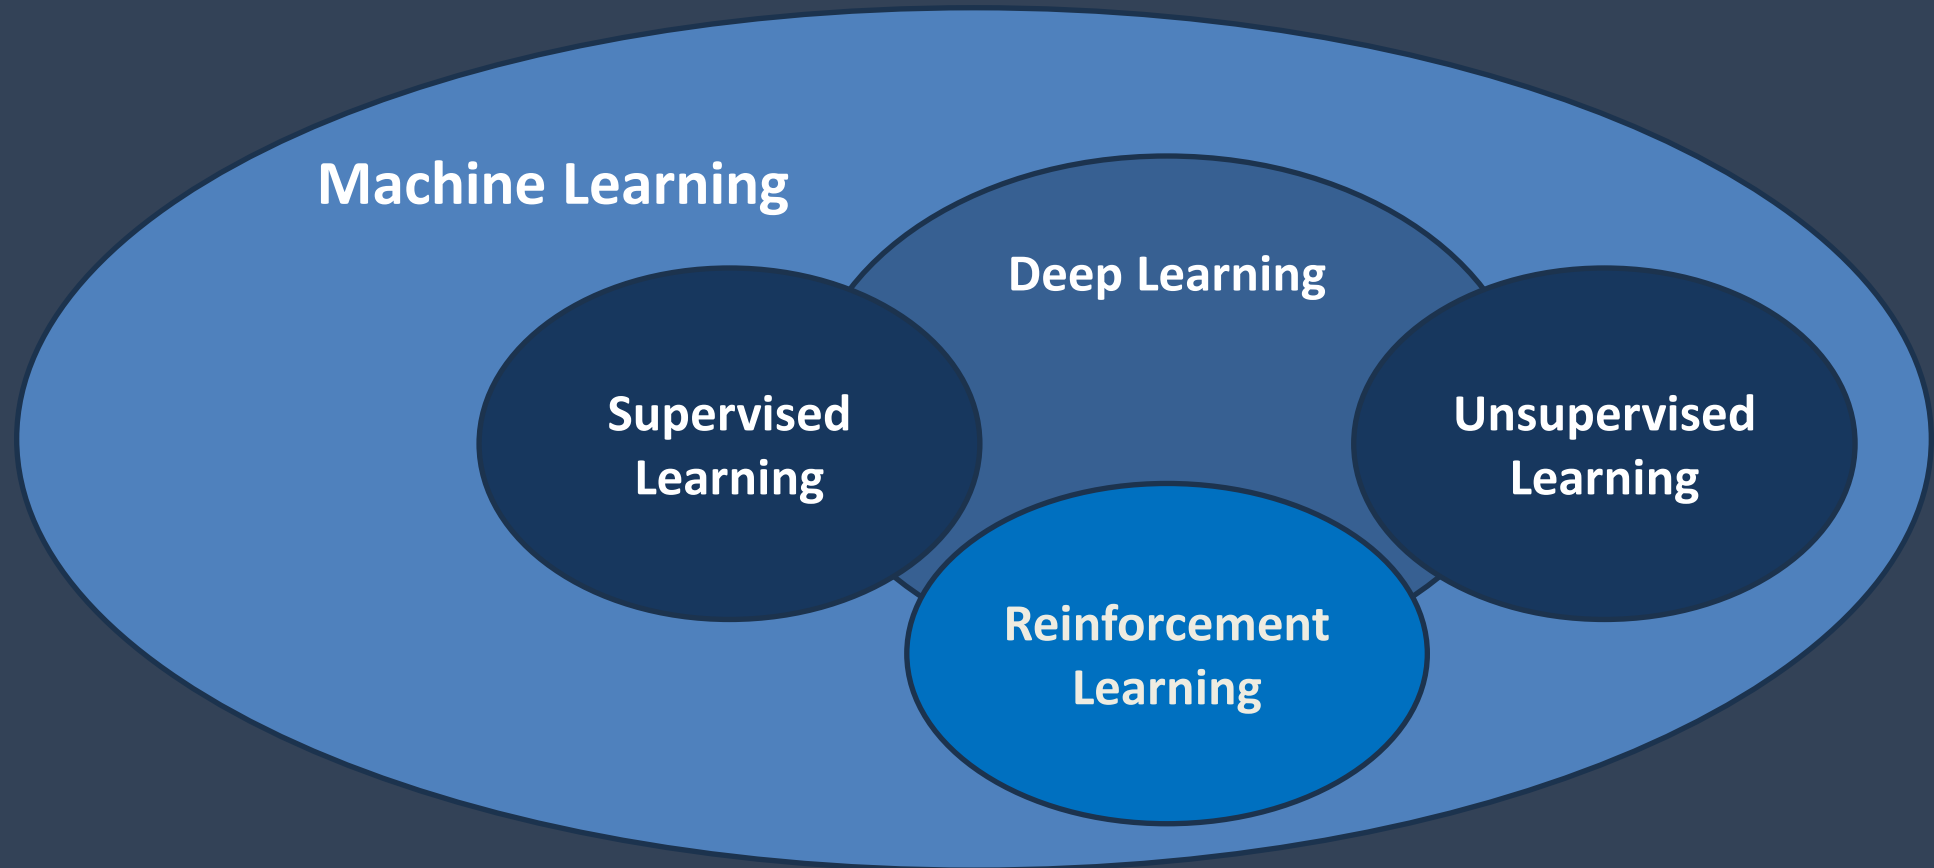
Professor: Li-Der Chou  
Presenter: Chih-Jou Tai

# Outline



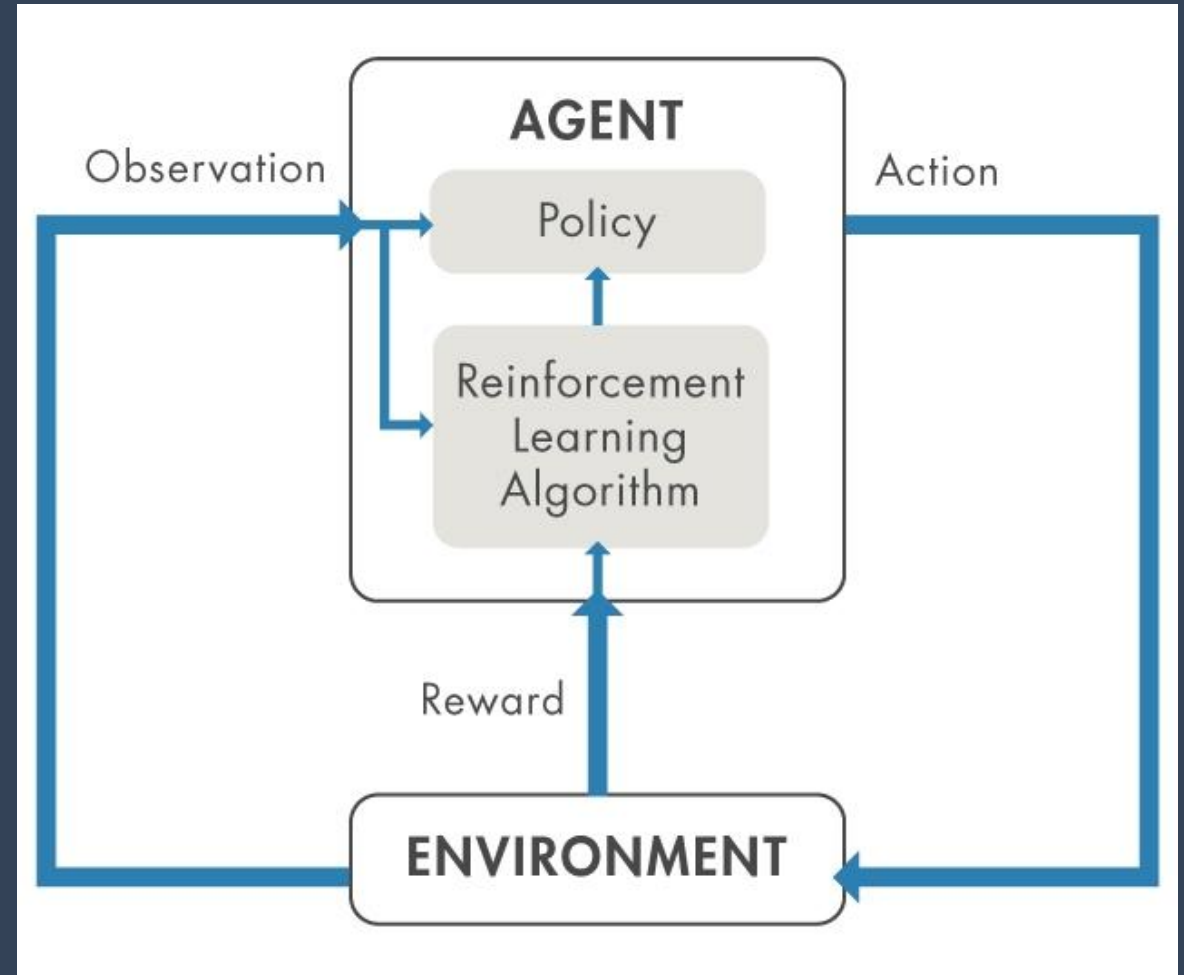
- Introduction
- Methodology
- Applications
- Pros and Cons
- References

- Reinforcement learning is a subfield within machine learning



# Introduction

- It involves the computer learning to perform a task correctly through continuous interaction with a dynamic environment, using trial and error.



# Outline



- Introduction
- Methodology
- Applications
- Pros and Cons
- References

# Methodology



1. Environment Setup
2. Reward Definition
3. Agent Creation
4. Agent Training and Validation
5. Policy Deployment

# Methodology-Environment Setup



- Define an environment in which the agent can learn, including specifying the interface between the agent and the environment.
- **Simulation Environment**: A virtual environment simulated by a computer, used for safe, controllable, and cost-effective experimentation and testing.
- **Experimental Environment**: An experimental environment is a physical environment in the real world, providing a more realistic context, but with lower control and safety and higher costs.

# Methodology-Reward Definition



- Define the agent's reward signals and explain how they are calculated, including factors like performance **metrics**, **costs**, **risks**, and more.
- **Discounting**: Discounting is about valuing immediate rewards more than delayed rewards in reinforcement learning by using a discount factor (often denoted as  $\gamma$ ) between 0 and 1.
- **Sparsity**: Sparsity in reinforcement learning means that rewards are rare or scarce, making the agent explore and figure out how to obtain them effectively.



- Selecting a policy representation method and choosing an appropriate training algorithm.
- **Policy Representation:** This step entails deciding how to represent the agent's policy. Common methods include using neural networks or lookup tables.
- **Training Algorithm:** Modern reinforcement learning often uses neural networks or Q-Learning for complex tasks with large state and action spaces. It determines how the agent adjusts its policy during learning.

- Configure training options and train the agent to adapt its policy. Validation of the trained policy is typically performed through simulations.
- **Training Configuration:** Including parameters like learning rates and exploration strategies.
- **Policy Validation:** After training, it's essential to assess the agent's policy to ensure it aligns with your goals. Validation is typically performed through simulations or real-world experiments.

# Methodology-Policy Deployment



- Determine how the trained policy will be deployed.
- **Programming Languages:** Common choices include C/C++, CUDA, Python, or specialized languages like TensorFlow or PyTorch.
- **Deployment Environment:** It can be embedded in a robot, integrated into software, or employed in a simulation.
- **Testing and Optimization:** Fine-tune and optimize it as needed to ensure it performs effectively in real-world scenarios.
- **Monitoring and Maintenance:** Monitor the policy's performance and be prepared to make adjustments or updates as necessary.

# Outline




- Introduction
- Methodology
- Applications
- Pros and Cons
- References

# Applications

- **Autonomous Driving**

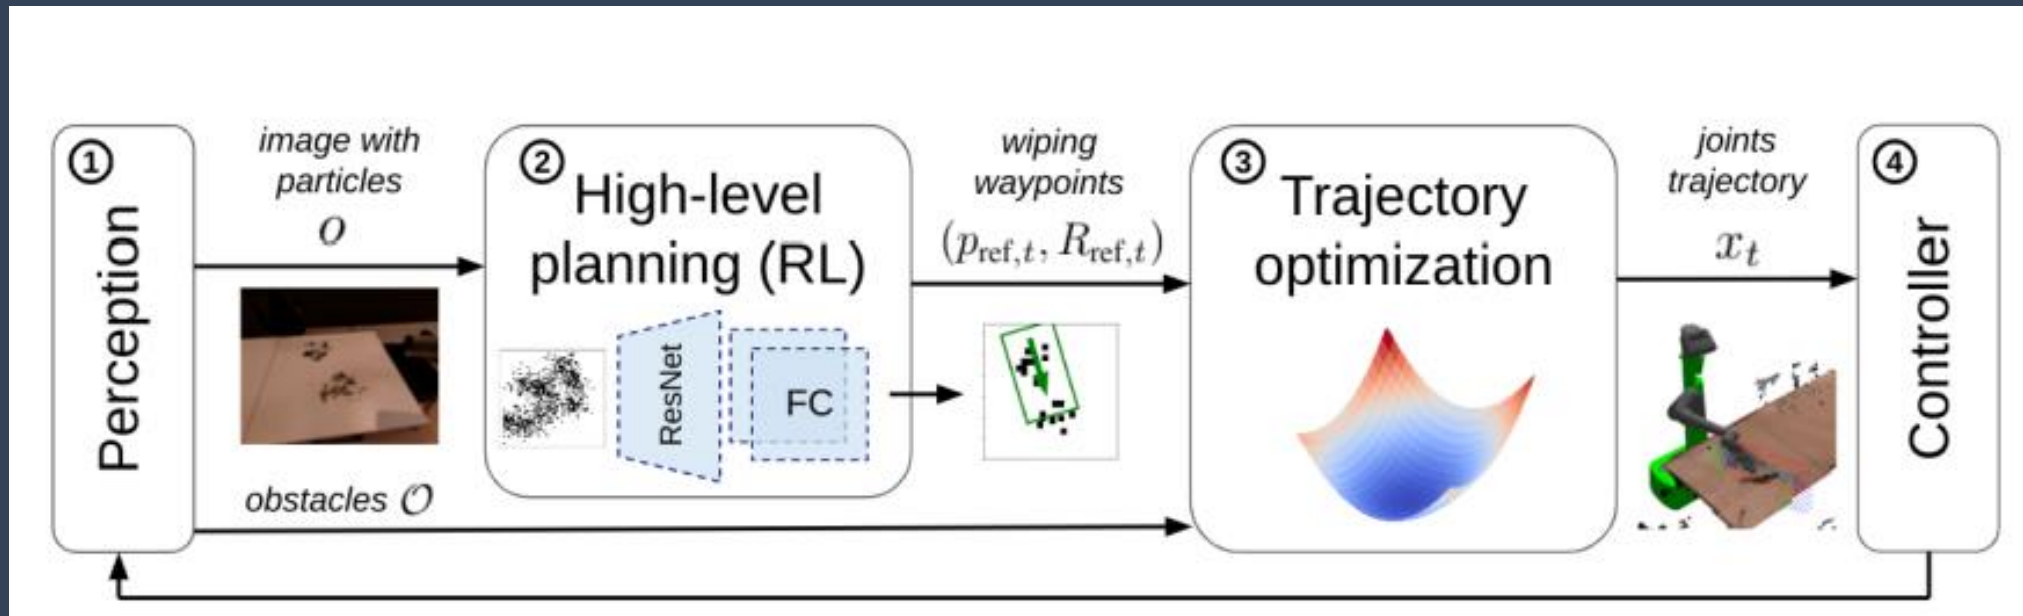
- In autonomous driving, the agent learns to navigate dynamic environments using sensors, policies driven by neural networks, and iterative reinforcement learning with rewards.



# Applications

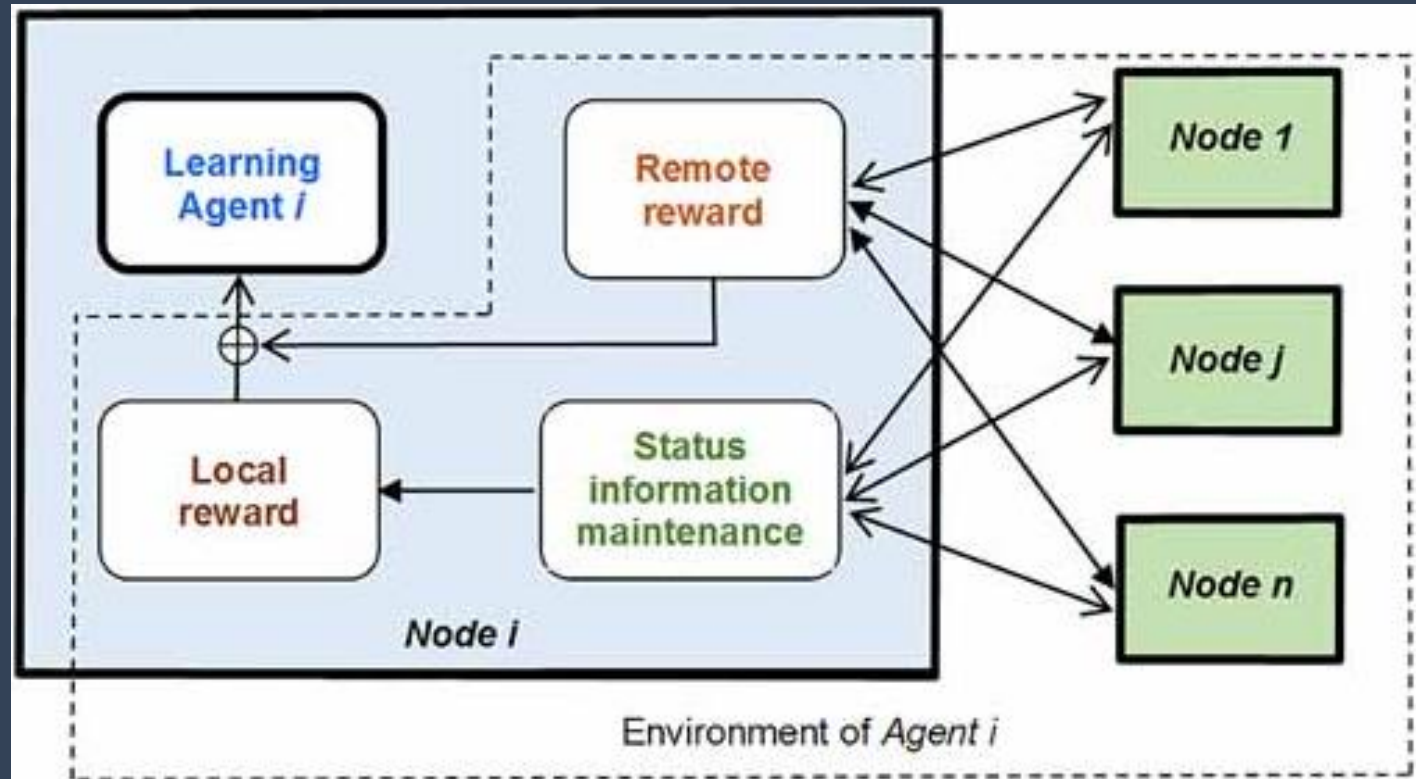
- **Robotic Table Wiping**

- This robot autonomously cleans tables by wiping away spills and crumbs using a combination of vision-based reinforcement learning and trajectory optimization for real-world deployment.



# Applications

- **Network Packet Routing**
- Agents (routers or switches) learn how to choose the optimal path to minimize packet transmission time or minimize network congestion.



# Outline



- Introduction
- Methodology
- Applications
- Pros and Cons
- References



# Pros



- **Versatility**: Applies to various tasks and domains, offering a versatile problem-solving approach.
- **Autonomy**: Agents learn and decide independently once the environment and algorithms are set up.
- **Adaptability**: Agents refine strategies through experience, enhancing performance over time.

# Cons

- **Sample Inefficiency**: Often demands substantial training data, leading to time and resource intensiveness.
- **Complex Problem Setup**: Designing problems correctly can be challenging, requiring multiple design iterations.
- **Black-Box Nature**: Trained neural network policies can be hard to interpret, affecting transparency.

# Outline



- Introduction
- Methodology
- Applications
- Pros and Cons
- References

# References



- 強化學習 (Reinforcement Learning)：入門指南  
[https://www.terasoft.com.tw/support/tech\\_articles/reinforcement\\_learning\\_a\\_brief\\_guide.asp](https://www.terasoft.com.tw/support/tech_articles/reinforcement_learning_a_brief_guide.asp)
- 強化學習訓練打掃機器人  
<https://arxiv.org/abs/2210.10865>
- Reinforcement Learning Based Routing in Networks  
<https://ieeexplore.ieee.org/document/8701570>

Q & A

